

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

Discrete Probability Distribution: Binomial

- A T-shirt manufacturing company regularly conducts quality checks at specified period on the products it manufactures. Their empirical data of manufacturing report says, 4% of their products were defective. Suppose a random sample of 250 T-shirts were picked from the manufacturing unit, what is the probability that,
 - a) None of the T-shirts are defective
 - b) 5 or fewer T-shirts are defective
 - c) 20 or more T-shirts are defective
 - d) What is the mean value of this Binomial Distribution
 - e) Plot the Binomial Distribution

```
In [4]: import scipy.stats as stats
```

```
In [ ]: #a) we need to calculate  $P(x=0) = n C x * P^x * Q^{(n-x)}$ 
#Here  $n=250$ ,  $x=0$ ,  $P=\text{defective \%} = 4\% = 0.04$ ,  $Q = \text{Non-defective \%} = 96\% = 0.96$ 
n=250
p=0.04
stats.binom.pmf(0,n,p)
```

```
Out[61]: 3.696649374485171e-05
```

```
In [ ]: #b) 5 or fewer =>  $P(x \leq 5) = P(x=0) + P(x=1) + P(x=2) + P(x=3) + P(x=4) + P(x=5)$ 
stats.binom.cdf(5,n,p)
```

```
Out[56]: 0.06329328366342199
```

```
In [ ]: #c) 20 or more =>  $P(x \geq 20) = P(x=20) + \dots + \text{all cumulative probabilities in right hand side}$ 
1-stats.binom.cdf(19,n,p)
```

Out[57]: 0.002814404938111137

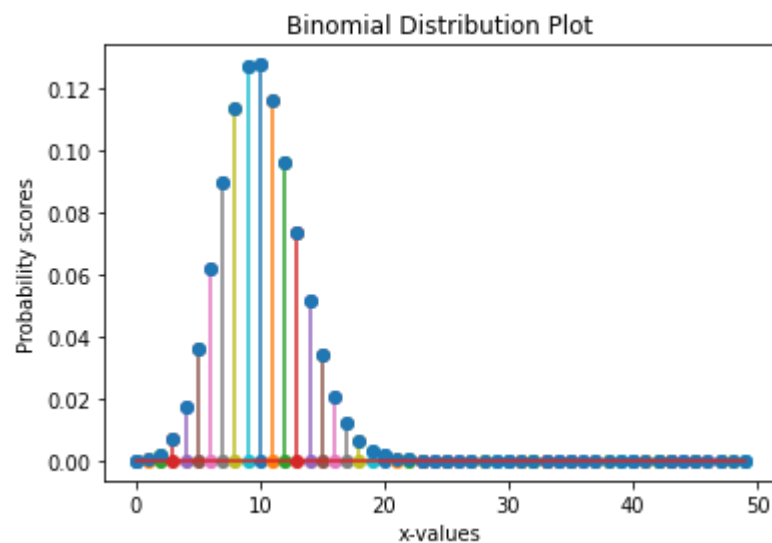
```
In [ ]: #d) Mean value of Binomial distribution =  $n \cdot p$ 
#Here, mean implies 4% of 250
0.04*250
```

Out[58]: 10.0

```
In [ ]: stats.binom.pmf(10,n,p)
```

Out[63]: 0.12768812205208388

```
In [ ]: #e) to generate the plot, we need to define x-axis
x=np.arange(0,50)
binom_dist=stats.binom.pmf(x,n,p)
plt.stem(x,binom_dist,'o-')
plt.xlabel('x-values')
plt.ylabel('Probability scores')
plt.title('Binomial Distribution Plot')
plt.show()
```



```
In [ ]: x.dot(binom_dist) #this is analogous to the notion summation of x.p(x)
```

```
Out[64]: 9.999999999999698
```

```
In [ ]: np.argmax(binom_dist), np.max(binom_dist)
```

```
Out[19]: (10, 0.12768812205208388)
```

```
In [ ]: sd=np.sqrt(((x-10)**2).dot(binom_dist)) #more generic formula  
sd
```

```
Out[65]: 3.0983866769658617
```

```
In [ ]: np.sqrt(n*p*(1-p))
```

```
Out[66]: 3.0983866769659336
```

- We could notice, the peak probability score of 12.76% at the mean value (10 T-shirts)

Discrete Probability Distribution: Poisson

A customer care unit of a fiber optic broadband service provider receives on an average 30 emails on a day regarding complaint in their service. To strategically plan the resource of service engineers, the following calculations are performed

-
- What is the probability of receiving 100 complaints in a given day?
 - What is the probability of receiving 20 or less complaints in a given day?
 - What is the probability of receiving 35 or more complaints in a given day?
 - Plot the poisson distribution

```
In [ ]: #a) we need to calculate  $P(x=100)$   
#Here Lambda = 30 emails (average)  
lamb=30  
stats.poisson.pmf(100,lamb)
```

Out[68]: 5.1675818018385396e-24


```
In [ ]: #b)  $P(x \leq 20)$   
stats.poisson.cdf(20,lamb)
```

Out[69]: 0.035284618454228846

```
In [ ]: #c)  $P(x \geq 35)$   
1-stats.poisson.cdf(34,lamb)
```

Out[70]: 0.20269167451688286

```
In [ ]: x=np.arange(0,60)
        pois_dist=stats.poisson.pmf(x,lamb)
        plt.stem(x,pois_dist,'o-')
        plt.xlabel('x-values')
        plt.ylabel('Probability scores')
        plt.title('Poisson Distribution Plot')
        plt.show()
```



```
In [ ]: np.argmax(pois_dist)
```

Out[73]: 29

```
In [ ]: pois_dist[29]
```

Out[74]: 0.07263452647159181

```
In [ ]: pois_dist[30]
```

Out[75]: 0.07263452647159181

- We can notice the highest probability centered around the mean (ie) 30 complaints

```
In [ ]: np.sqrt(x.dot(pois_dist)) #cross checking the mean value from histogram perspective x.P(x)
```

```
Out[80]: 5.477220430202007
```

```
In [ ]: np.sqrt(((x-30)**2).dot(pois_dist))
```

```
Out[79]: 5.477144707167058
```

#Normal Distribution | Value of the random variable | Area under the normal distribution (CDF) | | ----- | -----
 -----| | $\mu - \sigma \leq x \leq \mu + \sigma$ | 0.6828 | | $\mu - 2\sigma \leq x \leq \mu + 2\sigma$ | 0.9545 | | $\mu - 3\sigma \leq x \leq \mu + 3\sigma$ | 0.9974 |

- Based on last 5 year statistics, an app based food delivery company declared its statistics at Chennai, saying the average delivery time to any part of Chennai is 90 min, with a standard deviation of 20 min. Compute the following,
- a) What proportion of delivery were done less than 60 min?
- b) What proportion of delivery were done more than 100 min?
- c) What proportion of delivery were done between 30min to 60min?

```
In [ ]: #a)
stats.norm.cdf(60,90,20)
#answer is 6.68%
#.cdf by default will count the left hand
```

```
Out[5]: 0.06680720126885807
```

```
In [5]: #or alternatively we can pass
stats.norm.cdf(-1.5)
#we can convert the value into Z score and pass the value.
```

```
Out[5]: 0.06680720126885807
```

```
In [9]: 1-stats.norm.cdf(60,50,20)
```

```
Out[9]: 0.3085375387259869
```

```
In [ ]: (60-90)/20
```

```
Out[8]: -1.5
```

```
In [6]: #OR alternatively we can pass the  
stats.norm.cdf(-1.0)
```

```
Out[6]: 0.15865525393145707
```

```
In [10]: #b)  
1-stats.norm.cdf(100,90,20)
```

```
Out[10]: 0.3085375387259869
```

```
In [11]: #c)  
stats.norm.cdf((60-90)/20)-stats.norm.cdf((30-90)/20)
```

```
Out[11]: 0.06545730323722798
```

- Inverse Survival function gives the value given a probability

```
In [ ]: #What is the cutoff value for 5% area in standard normal curve?  
#2.5% either side of the curve  
2.5/100
```

```
Out[19]: 0.025
```

```
In [12]: stats.norm.isf(0.025) #95% area
```

```
Out[12]: 1.9599639845400545
```

```
In [ ]: stats.norm.isf(0.005) #99% area
```

```
Out[10]: 2.575829303548901
```

```
In [13]: stats.norm.isf(0.05) #90% area
```


```
Out[13]: 1.6448536269514729
```

```
In [15]: import random  
pop_sal=np.random.randn(50000)  
pop_sal  
#simulating 50000 of the population
```

```
Out[15]: array([-0.6889953 ,  0.46889536,  0.23472991, ..., -0.57287622,  
                0.41898068,  0.0233784 ])
```

```
In [16]: pop_sal_df=pd.DataFrame(pop_sal)  
pop_sal_df.hist()
```

```
Out[16]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fcad12d5358>]],  
              dtype=object)
```




```
In [19]: #convert all 50000 samples in real INR units (assume pop_mean(true_mean=6.78LPA,sigma=40k))  
sal=pop_sal*40000+678000  
sal[:50]#just for display
```

```
Out[19]: array([650440.18816621, 696755.81443549, 687389.19627987, 577756.73698671,  
680252.23970453, 667314.90670656, 702169.03126286, 604371.99263918,  
714377.71707805, 646829.52412998, 683683.72694913, 691232.142362 ,  
721821.70520121, 656616.62369856, 732765.70370266, 665166.57252265,  
668233.9903018 , 611905.56246847, 661503.73512871, 665964.43539245,  
629421.84009142, 664270.24106678, 695880.24398613, 721860.70182764,  
686795.62710355, 733167.93087388, 686468.0575932 , 666492.27814401,  
678659.88698166, 632251.08709472, 652891.3067268 , 663805.47579432,  
702513.78957679, 765021.64205576, 738109.13540897, 737951.70536068,  
643509.42933873, 696239.88287976, 715224.77919337, 696034.67589335,  
685827.22912836, 714394.02661737, 724932.3531023 , 722378.44347268,  
692728.57689258, 702100.39727511, 681792.80248471, 628897.39377278,  
703697.28426036, 621250.90480723])
```

```
In [20]: np.mean(sal)
```

```
Out[20]: 678089.8193874795
```


```
In [21]: #Method-1 (as per the theoretical definition of CLT)  
trial=[]  
for itr in np.arange(100):  
    samp=[]  
    for val in np.arange(500):  
        samp.append(np.random.choice(sal))  
    trial.append(np.mean(samp))
```

```
In [22]: len(trial)
```

```
Out[22]: 100
```

```
In [25]: trial_df=pd.DataFrame(trial)
trial_df.hist()
```

```
Out[25]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fcad11dbd68>]],
      dtype=object)
```



```
In [24]: np.mean(trial)
```

```
Out[24]: 678166.1314576634
```

```
In [ ]: np.mean(trial),np.std(trial,ddof=1)
```

```
Out[30]: (678171.6527113924, 1693.8784830311258)
```

```
In [26]: #95% CI range estimate of pop_mean is given by  
(678171.6527113924-1.96*1693.8784830311258,678171.6527113924+1.96*1693.8784830311258)
```

```
Out[26]: (674851.6508846513, 681491.6545381334)
```

```
In [27]: #68% CI range estimate of pop_mean is given by  
(678171.6527113924-1*1693.8784830311258,678171.6527113924+1*1693.8784830311258)
```

```
Out[27]: (676477.7742283612, 679865.5311944236)
```

```
In [ ]: trial_df=pd.DataFrame(trial)  
trial_df.std()
```

```
Out[24]: 0    1693.878483  
dtype: float64
```

```
In [ ]: np.std(trial,ddof=1)
```

```
Out[26]: 1693.8784830311258
```

```
In [ ]: #(95% range of this sal)  
678171.6527113924-1.96*1693.8784830311258
```

```
Out[31]: 674851.6508846513
```

```
In [ ]: 678171.6527113924+1.96*1693.8784830311258
```

```
Out[32]: 681491.6545381334
```

```
In [ ]: #[6.74Lacs to 6.81Lacs]
```

```
In [ ]: #Method-2 (practical approach with approx)  
samp=[]  
for val in np.arange(500):  
    samp.append(np.random.choice(sal))
```

```
In [ ]: np.mean(samp),np.std(samp,ddof=1)
```

```
Out[36]: (681015.0281048274, 41018.99285935366)
```

```
In [ ]: #Standard Error of Mean in practical approach  
sem=np.std(samp,ddof=1)/np.sqrt(500)  
sem
```

```
Out[37]: 1834.425128041865
```

```
In [ ]: 681015.0281048274-1.96*sem
```

```
Out[38]: 677419.5548538653
```

```
In [ ]: 681015.0281048274+1.96*sem
```

```
Out[39]: 684610.5013557895
```

```
In [ ]: #De-moivre's approximation of Discrete Distribution with Standard Normal Distribution  
#Generate Binomial distribution with p=0.5, n=25 and calculate P(x<=14)  
stats.binom.cdf(14,25,0.5)
```

```
Out[32]: 0.7878218889236449
```

```
In [ ]: #If x is a normal random variable,mean=n*p, sd=npq  
x_bar=25*.5  
x_bar
```

```
Out[34]: 12.5
```

```
In [ ]: x_var=25*0.5*0.5  
x_sd=np.sqrt(x_sd)
```

```
Out[36]: 2.5
```

```
In [ ]: (14-12.5)/2.5
```

```
Out[37]: 0.6
```

*In discrete distribution, the histogram bars are centered on the numbers. This means $P(x \leq 14)$ in Discrete distribution is actually the area under the bars less than $x=14.5$, we need to account for that extra 0.5, while calculating the same area in continuous distribution

```
In [ ]: (14.5-12.5)/2.5
```

```
Out[42]: 0.8
```

```
In [ ]: #we need to calculate the prob P(x<=14) which equals to P(z<=0.6)  
stats.norm.cdf(0.6)
```

```
Out[41]: 0.7257468822499265
```

```
In [ ]: stats.norm.cdf(0.8)
```

```
Out[43]: 0.7881446014166034
```

- Now we have approximated Discrete distribution using Normal Distribution [$0.7878 = 0.7881$ (approx)]

```
In [ ]:
```