

```
In [1]: import numpy as np
import pandas as pd
import scipy.stats as stats
from scipy.stats import ttest_ind
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: import random
vol=np.random.randn(400)
#randn will generate values in standard normal scale
vol[:50]
```

```
Out[3]: array([-3.26328375e-01, -3.75551486e-01, -1.22148074e+00,  4.75539527e-01,
  2.07935050e+00, -1.93233821e+00,  3.03882745e-01,  4.65615859e-01,
 -2.92467949e-01, -2.77921169e+00, -3.71255142e-02,  4.31287181e-01,
  2.54679137e+00,  2.47497081e-01, -6.81927992e-01,  2.20881442e+00,
 -1.09452715e+00, -3.14261720e-01, -1.09431824e+00,  1.46915824e+00,
 -5.21034048e-01, -1.16939168e-01, -1.25526832e+00, -1.87880152e+00,
 -3.29281416e-01, -2.34937478e-01, -3.47567045e-01,  6.20307171e-01,
  2.01675941e-01,  1.45150895e-03,  5.43461913e-01, -1.20237380e-01,
  1.77981442e+00, -7.93108750e-01,  7.81308358e-01,  1.31609768e+00,
  9.05585817e-01,  7.83014183e-01,  3.05764448e-01, -2.63191548e+00,
 -2.29322771e+00,  5.15631126e-01, -2.24448697e-01,  7.62143653e-01,
  2.60881292e-01,  1.95498272e+00,  9.36165787e-01, -1.96664580e-01,
  1.82994294e+00,  5.20215195e-01])
```

```
In [4]: samp_vol=vol*1.5+298.56  
        samp_vol[:50]
```

```
Out[4]: array([298.07050744, 297.99667277, 296.7277789 , 299.27330929,  
              301.67902575, 295.66149269, 299.01582412, 299.25842379,  
              298.12129808, 294.39118247, 298.50431173, 299.20693077,  
              302.38018705, 298.93124562, 297.53710801, 301.87322164,  
              296.91820928, 298.08860742, 296.91852263, 300.76373735,  
              297.77844893, 298.38459125, 296.67709752, 295.74179772,  
              298.06607788, 298.20759378, 298.03864943, 299.49046076,  
              298.86251391, 298.56217726, 299.37519287, 298.37964393,  
              301.22972163, 297.37033688, 299.73196254, 300.53414652,  
              299.91837873, 299.73452127, 299.01864667, 294.61212678,  
              295.12015844, 299.33344669, 298.22332695, 299.70321548,  
              298.95132194, 301.49247407, 299.96424868, 298.26500313,  
              301.30491442, 299.34032279])
```

```
In [5]: #To verify fill volume of softdrink follows 300ml spec  
#Random 400 samples were collected from the factory  
#The statistics is found to be  
x_bar=np.mean(samp_vol)  
s=np.std(samp_vol,ddof=1)  
x_bar,s
```

```
Out[5]: (298.51331860969555, 1.549862841816424)
```

One sample t-Test (two tailed)

- Ho: pop\_mean=300 ml
- Ha: pop\_mean != 300 ml

```
In [6]: t_stat=(x_bar-300)/(s/np.sqrt(400))  
        t_stat
```

```
Out[6]: -19.184683317680918
```

```
In [ ]: # we can verify the score with the built in function
```

```
In [7]: from scipy.stats import ttest_1samp
```

```
In [9]: stats.norm.isf(0.05)
```

```
Out[9]: 1.6448536269514729
```

```
In [8]: ttest_1samp(samp_vol,300)
```

```
Out[8]: Ttest_1sampResult(statistic=-19.184683317680918, pvalue=1.3520986667089086e-58)
```

since p-val < 0.05, we reject  $H_0$ , which implies  $H_a$  holds good, (ie) sample doesn't represent 300 ml

```
In [40]: #Courier Company Example (One-tail test)
#n=50
dtime=np.random.randn(50)
samp_dtime=dtime*0.6+2.88
```

```
In [41]: x_bar=np.mean(samp_dtime)
s=np.std(samp_dtime,ddof=1)
x_bar,s
```

```
Out[41]: (2.9690997558267496, 0.6322643868481541)
```

```
In [48]: #compute 95% CI range for x_bar
x_bar-1.96*(s/np.sqrt(50))
```

```
Out[48]: 2.7938450091330607
```

```
In [49]: x_bar+1.96*(s/np.sqrt(50))
```

```
Out[49]: 3.1443545025204385
```

```
In [50]: t_stat=(x_bar-3)/(s/np.sqrt(50))  
t_stat
```

```
Out[50]: -0.3455796759983094
```

```
In [51]: ttest_1samp(samp_dtime,3)
```

```
Out[51]: Ttest_1sampResult(statistic=-0.3455796759983094, pvalue=0.7311378981078815)
```

```
In [ ]: s
```

```
Out[16]: 0.5627626509636148
```

```
In [ ]: t_stat=(x_bar-3)/(s/np.sqrt(500))  
t_stat
```

```
Out[31]: -7.5972214726259075
```

```
In [ ]: stats.norm.isf(0.05)
```

```
Out[13]: 1.6448536269514729
```

```
In [ ]: np.mean(samp_dtime),np.std(samp_dtime,ddof=1)
```

```
Out[30]: (2.791987829101012, 0.6122361391640108)
```

```
In [ ]: ttest_1samp(samp_dtime,3)
```

```
Out[28]: Ttest_1sampResult(statistic=-7.597221472625907, pvalue=1.505809226022567e-13)
```

since,  $p\text{-val} < 0.05$ , we reject  $H_0$ , which implies  $H_a$  holds good, ie (courier company claim is correct)

```
In [ ]: x_bar-1.96*(s/np.sqrt(50))
```

```
Out[25]: 2.6222846294303976
```

```
In [ ]: x_bar+1.96*(s/np.sqrt(50))
```

```
Out[26]: 2.9616910287716265
```

Soyabean yield example (one-tailed-Right tailed)

- $H_a$ : pop\_mean > 520
- $H_o$ : pop\_mean ≤ 520

```
In [52]: yd=np.random.randn(400)  
samp_yd=yd*124+573
```

```
In [55]: x_bar=np.mean(samp_yd)  
s=np.std(samp_yd,ddof=1)  
x_bar,s
```

```
Out[55]: (572.8865052001379, 123.7503003973732)
```

```
In [56]: t_stat=(x_bar-520)/(s/np.sqrt(400))  
t_stat
```

```
Out[56]: 8.547293223582429
```

```
In [57]: ttest_1samp(samp_yd,520)
```

```
Out[57]: Ttest_1sampResult(statistic=8.547293223582429, pvalue=2.7099868415520073e-16)
```

```
In [60]: #HYD DSE Data  
dse_age_hyd=np.random.randn(800)  
DH_age=dse_age_hyd*1.8+24.3  
np.mean(DH_age),np.std(DH_age,ddof=1)
```

```
Out[60]: (24.331701058427495, 1.7577885499753965)
```

```
In [61]: #BLR DSE Data
dse_age_blr=np.random.randn(1000)
DB_age=dse_age_blr*2.3+26.5
np.mean(DB_age),np.std(DB_age,ddof=1)
```

```
Out[61]: (26.505777478090163, 2.321083438150285)
```

```
In [63]: ttest_ind(DB_age,DH_age)
```

```
Out[63]: Ttest_indResult(statistic=21.93417098536577, pvalue=1.088058828703349e-94)
```

```
In [64]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## Two Sample t Test

- $H_0: \mu_1 = \mu_2$
- $H_a: \mu_1 \neq \mu_2$

$$\text{Test statistic } T = \frac{\overline{X_1} - \overline{X_2}}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

- where  $n_1$  and  $n_2$  are the sample sizes and  $X_1$  and  $X_2$  are the sample means
- $S_1^2$  and  $S_2^2$  are sample variances


```
In [67]: A=pd.read_table('/content/drive/My Drive/Statistics Mahesh Anand/HR.txt',index_col=0)  
A.head()
```

Out[67]:

	Attrition	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	Employ
--	-----------	-----	----------------	-----------	------------	------------------	-----------	----------------	---------------	--------

Individual

Ind1	Yes	41	1	1102	1	1	2	1	1	
Ind2	No	49	2	279	2	8	1	1	1	
Ind3	Yes	37	1	1373	2	2	2	6	1	
Ind4	No	33	2	1392	2	3	4	1	1	
Ind5	No	27	1	591	2	2	1	2	1	



```
In [68]: A.loc['Ind400']
```

```
Out[68]: Attrition                No
Age                31
BusinessTravel     1
DailyRate         329
Department        2
DistanceFromHome   1
Education          2
EducationField     1
EmployeeCount      1
EmployeeNumber     530
EnvironmentSatisfaction  4
Gender            1
HourlyRate        98
JobInvolvement     2
JobLevel          1
JobRole           3
JobSatisfaction    1
MaritalStatus      2
MonthlyIncome     2218
MonthlyRate       16193
NumCompaniesWorked 1
OverTime          2
PercentSalaryHike  12
PerformanceRating  3
RelationshipSatisfaction  3
StandardHours     80
StockOptionLevel   1
TotalWorkingYears  4
TrainingTimesLastYear  3
WorkLifeBalance    3
YearsAtCompany     4
YearsInCurrentRole  2
YearsSinceLastPromotion  3
YearsWithCurrManager  2
Name: Ind400, dtype: object
```



```
In [69]: A.shape
```

```
Out[69]: (1470, 34)
```

```
In [83]: A['Gender'].value_counts()
```

```
Out[83]: 1    882
         2    588
         Name: Gender, dtype: int64
```

- Verify the avg sal of employee who left the org = avg sal of emp who currently working in the org
- Ho: pop\_mean\_sal (left)=pop\_mean\_sal (currently working)
- Ha: not equal

```
In [77]: df_yes=A[A['Attrition']=='Yes']
         df_no=A[A['Attrition']=='No']
```

```
In [84]: df=A.groupby('Gender')
         df1=df.get_group(1)
         df2=df.get_group(2)
         df1.shape,df2.shape
```

```
Out[84]: ((882, 34), (588, 34))
```

```
In [78]: ttest_ind(df_yes['MonthlyIncome'],df_no['MonthlyIncome'])
```

```
Out[78]: Ttest_indResult(statistic=-6.203935765608938, pvalue=7.14736398535381e-10)
```

```
In [85]: ttest_ind(df1['MonthlyIncome'],df2['MonthlyIncome'])
```

```
Out[85]: Ttest_indResult(statistic=-1.2212617308870655, pvalue=0.22218303455087898)
```

```
In [86]: df1['MonthlyIncome'].mean()
```

```
Out[86]: 6380.507936507936
```

```
In [87]: df2['MonthlyIncome'].mean()
```

```
Out[87]: 6686.566326530612
```

```
In [79]: df_yes['MonthlyIncome'].mean()
```

```
Out[79]: 4787.0928270042195
```

```
In [80]: df_no['MonthlyIncome'].mean()
```

```
Out[80]: 6832.739659367397
```

```
In [ ]: A.columns
```

```
Out[37]: Index(['Attrition', 'Age', 'BusinessTravel', 'DailyRate', 'Department',  
              'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',  
              'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',  
              'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',  
              'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',  
              'OverTime', 'PercentSalaryHike', 'PerformanceRating',  
              'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',  
              'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',  
              'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',  
              'YearsWithCurrManager'],  
              dtype='object')
```

```
In [ ]: A.shape
```

```
Out[38]: (1470, 34)
```

```
In [ ]: A['Attrition'].value_counts()
```

```
Out[39]: No      1233  
        Yes      237  
        Name: Attrition, dtype: int64
```

```
In [ ]: df_yes=A[A['Attrition']=='Yes']  
df_no=A[A['Attrition']=='No']  
df_yes.shape,df_no.shape
```

```
Out[68]: ((237, 34), (1233, 34))
```

```
In [ ]: #OR  
DF=A.groupby('Attrition')  
df1=DF.get_group('Yes')  
df2=DF.get_group('No')  
df1.shape,df2.shape
```

```
Out[67]: ((237, 34), (1233, 34))
```

```
In [ ]: sal_yes=df_yes['MonthlyIncome']  
sal_no=df_no['MonthlyIncome']
```

```
In [ ]: np.mean(sal_yes),np.mean(sal_no)
```

```
Out[72]: (4787.0928270042195, 6832.739659367397)
```

```
In [ ]: ttest_ind(sal_no,sal_yes)
```

```
Out[71]: Ttest_indResult(statistic=6.203935765608938, pvalue=7.14736398535381e-10)
```

Verify the average sal of employees who left the organisation is same as avg sal of emp currently working in organisation

```
In [ ]: sal_yes=A['MonthlyIncome'][A['Attrition']=='Yes']  
sal_no=A['MonthlyIncome'][A['Attrition']=='No']
```

```
In [ ]: ttest_ind(sal_yes,sal_no)
```

```
Out[42]: Ttest_indResult(statistic=-6.203935765608938, pvalue=7.14736398535381e-10)
```

p-val<5%(0.05), which implies,we are rejecting Ho, ie,avg sal is not equal (significant) difference

```
In [ ]: np.mean(sal_yes),np.mean(sal_no)
```

```
Out[43]: (4787.0928270042195, 6832.739659367397)
```

Since,  $p\text{-val} < 0.05$  (5%) which falls in the acceptance zone of  $H_a$ . Hence we reject  $H_0$  with 95% confidence. There is a statistical evidence to say the two means are significantly different

```
In [ ]: A.columns
```

```
Out[9]: Index(['Attrition', 'Age', 'BusinessTravel', 'DailyRate', 'Department',
              'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
              'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
              'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
              'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
              'OverTime', 'PercentSalaryHike', 'PerformanceRating',
              'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
              'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
              'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
              'YearsWithCurrManager'],
             dtype='object')
```

```
In [ ]: sal_male=A['MonthlyIncome'][A['Attrition']=='Yes']
        sal_female=A['MonthlyIncome'][A['Attrition']=='No']
```

```
In [ ]: A['Gender'].value_counts()
```

```
Out[10]: 1    882
         2    588
         Name: Gender, dtype: int64
```

```
In [ ]: sal_male=A['MonthlyIncome'][A['Gender']==1]
        sal_female=A['MonthlyIncome'][A['Gender']==2]
```

```
In [ ]: ttest_ind(sal_male,sal_female)
```

```
Out[45]: Ttest_indResult(statistic=-1.2212617308870655, pvalue=0.22218303455087898)
```

$p\text{-val} > 5\%$  (0.05) infact it is 22%, we have strong evidence to accept  $H_0$  ( $\text{Sal\_male\_emp} = \text{Sal\_female\_emp}$ )

```
In [ ]: np.mean(sal_male),np.mean(sal_female)
```

```
Out[16]: (6380.507936507936, 6686.566326530612)
```

Since p-value >0.05 (5%) it falls in acceptance zone of  $H_0$ , ie., average salary of male and female are nearly same.